

EE1301

Lab 8: Pointers and Dynamic Memory

The explicit memory management capability of C/C++ is one reason that the language remains popular nearly half a century after its creation. Using dynamic memory requires an understanding of *indirection* – the idea of manipulating a variable (*pointer variable*) that refers to another variable. We say that the pointer variable *points* to the the location in memory where another variable is stored. In other words, the pointer stores the address of the other variable. In this lab exercise we will explore dynamic memory management.

Warm-up

Complete the following paper/pencil exercises. First, work on them individually. Then, discuss them with a neighbor and make any needed corrections. Finally, discuss your results with your TA.

1) Pointer basics

- a) Declare two pointer variables of type `double` named `d_var` and `d_array`.
- b) Write C++ statements to dynamically create a double-precision floating-point variable and store its address in `d_var`. Also, dynamically create an array of 10 double-precision floating-point values and store the address of the array in `d_array`.
- c) Write C++ statements to request an input value for the variable that `d_var` points to. Get the value from the console and then display it.
- d) Write C++ statements to initialize all 10 of the `double` values in the dynamically allocated array to 1.0.
- e) Now, write C++ statements to de-allocate the memory referenced by `d_var` and `d_array`, using the `delete` operator.

2) Pointer management

- a) Show the output of the following code segment.

```
int a(1);
int b(2);
int *p1, *p2;
p1 = &a;
p2 = &b;
*p1 = *p2;
*p2 = 10;
cout << *p1 << ' ' << b << ' ' << a << endl;
```

- b) Describe the problem with the following function.

```
void foo()
{
    int *array = new int[100];
    for(int i = 0; i < 100; i++)
    {
        //do something
    }
}
```

3) Pointers with classes

a) A user-defined class named `Timer` has a constructor that takes two integer arguments to initialize `hour` and `minute` data members. Write a single C++ statement that creates an object of the `Timer` class using **dynamic memory allocation** and assigns the pointer returned by `new` to a pointer variable named `timePtr`. The statement should call the constructor with two parameters. Use values of 10 and 20 for `hour` and `minute`, respectively.

b) Write the definition for a function named `randArray` that takes a single integer argument, `n`, and returns a dynamically-allocated array of `n` pseudo-random integers. You may assume that the pseudo-random number generator has been previously declared and seeded (i.e. you do not need `srand(time(0))` or a `#include`).

c) Now, write C++ statements to call the `randArray` function to construct a 100-element array. Then, print the array values to the terminal (one per line) and delete the dynamically allocated array.

Stretch

1) Momentum

Momentum is defined as the product of an item's *mass* and its *velocity*. *Mass* is a scalar value, whereas *velocity* is expressed as a vector quantity with three components that describe the velocity in the x, y, and z directions. The product of the scalar *mass* and the *velocity* vector yields *momentum* as a vector quantity.

Write a function named `momentum` that accepts as arguments:

(i) a one-dimensional *velocity* array with three values of type `double` (i.e. a 3d vector) and

(ii) a *mass* of type `double`,

and returns a *dynamically-allocated* array representing the *momentum*. Note that *momentum* is determined by multiplying the scalar *mass* by each element of the vector *array*.

Test your `momentum` function by constructing a short `main` program that asks the user to input values for the velocity and mass from the console and then displays the momentum.

Workout

1) Arrays of Dynamic Items

Write a program to determine the average momentum of a collection of items with random velocities and masses. Do this using the following outline.

1. Construct a function named `randVec` that takes no arguments and returns a dynamically-allocated 3-element array of doubles. Each element in the array should be a randomly-generated value in the range `[-100.0, +100.0]`.
2. Using `randVec` and your `momentum` function from the previous part, generate *momentum* vectors for 1000 items, each of which has a random *velocity* (as described above) and a randomly-generated *mass* in the range `[1.0, 10.0]`. Save the *momentum* vectors using a suitable array of pointers.
3. Compute and display the average momentum vector of the items using a *for* loop. [Hint: the average should be calculated component by component.]

2) Letter Frequency

Write a function that takes a string as input and returns a counter for each letter in the string. For example, "my dog ate my homework" contains 3 m's, 3 o's, 2 e's, 2 y's and one each of d, g, a, t, h, w, r, and k.

Your function should take a single string argument and return a *dynamically allocated* array of 26 integers representing the count of each of the letters a – z, respectively. Your function should be case *insensitive*, i.e., count both 'A' and 'a' as an occurrence of the letter a. [Hint: Use the letter to integer conversion functions.] Do not count non-letter characters (i.e., spaces, punctuation, digits, etc.).

Write a program that requests a string from the user using `getline`, calls your letter frequency function, and prints out the frequency of each letter in the string. Do not print out letters that do not occur at least once.

Example:

```
Enter a string: my dog at my homework
Letter frequency
a  1
d  1
e  1
g  1
h  1
k  1
m  3
o  3
r  1
t  1
w  1
y  2
```

Challenge

1) Buffering

Real-world data from measuring devices or sensors is generally produced *asynchronously* relative to the processing routine(s). The data may appear in bursts that occur faster than they can be individually processed. An effective way to handle this situation is to employ a *buffer*. A *buffer* is simply an array into which a data-*generating* process writes data, and from which a separate data-*processing* routine subsequently retrieves the data when it is ready to do processing. In the short-term, if the data is being generated at a faster rate than it can be processed, the buffer fills up while the processing routine catches up. In the long-term (assuming the average processing rate exceeds the average data generation rate), the data processing will be able to remain ahead of or keep pace with the data-generating process.

There are a number of ways to implement buffering. One simple way is to employ a *double-buffering* technique in which a generating process allocates two buffers. At any given time, the data-generating process is filling one of the buffers while allowing the processing routine to read from the other buffer. Once the data-generating process fills one buffer, it provides the address of the filled buffer to the processing routine and allocates a new buffer to begin filling. The processing routine extracts the data from the filled buffer and disposes the buffer when finished, then awaits a pointer to the next "filled" buffer.

In this problem, you will simulate a basic *double-buffering* scheme using dynamically-allocated arrays of integers. You will need to construct a data *generation* function and a data *processing* function and then randomly call them from a simulation loop in the main program.

- 1) Construct a function: `double getProb()` that will return a pseudo-random value between 0.0 and 1.0.
- 2) Construct a function: `int* generateData(int* &inbuf, int &count)` that simulates asynchronous data generation by obtaining a random number between 0 and 9 (you do not need to use `getProb()` for this) and saving it in an input buffer. Your function should do the following:
 - Generate a random integer between 0 and 9 and insert it in the input buffer at the array location specified by the count argument. (Note that the buffer is passed as a pointer variable reference.)
 - Increase the count.
 - If the buffer is full, (a) return the address of the full buffer, (b) reset the count to zero, (c) allocate a new buffer, and (d) save its address in the pointer variable passed as the first argument.
 - If the buffer is not full, then return NULL.
- 3) Construct a function: `void processData(int* &outbuf, int &count, int &total)` to simulate "processing" of the asynchronous data. Do the following:
 - If the output buffer pointer is NULL, do nothing, i.e., just return.
 - Otherwise, obtain the buffer value at `[count]` and add it to `total`, then increase the count.
 - If all elements in the buffer have been exhausted, then reset the count to zero, delete the (dynamically-allocated) buffer and set the buffer pointer argument to NULL.
- 4) Finally, include the following code in a main simulation program that will call the `generateData` and `ProcessData` functions in a loop.

```
const int BUFSIZE=10;
const int ITERATIONS=50;

int main()
{
    int *fillbuffer = new int[BUFSIZE];
    int fillcnt=0;
    int *processbuffer = NULL;
    int processcnt=0;
    int tcount = 0;

    for(int i=0; i<ITERATIONS; i++)
    {
        int *temp;
        if(getProb() <= 0.40 )
        {
            temp = generateData(fillbuffer,fillcnt);
            if( temp != NULL ){
                processbuffer = temp;
            }
        }
        if(getProb() <= 0.60){
            processData(processbuffer,processcnt,tcount);
        }
        cout << fillcnt << '\t' << processcnt << endl;
    }
    cout << "Total value: " << tcount << endl;
    return 0;
}
```